

ON THE PITFALLS OF VERBALIZED CONFIDENCE PRIORS FOR CALIBRATING REASONING LLMs

Anonymous authors

Paper under double-blind review

ABSTRACT

The abstract paragraph should be indented 1/2 inch (3 picas) on both left and qaqq

1 INTRODUCTION

Reasoning large language models (LLMs) increasingly solve complex problems by generating a reasoning trajectory followed by the answer (Yang et al., 2025). Despite these advances, the answer accuracy alone does not ensure reliable uncertainty estimates: a model may produce plausible reasoning trajectory yet remain overconfident in an incorrect answer (Geng et al., 2024; Liu et al., 2025). In principle, confidence calibration requires the reported confidence to reflect the empirical probability that the answer is correct, which in turn enables applications such as hallucination detection (Wu et al., 2026) and adaptive test-time scaling (Yang et al., 2026).

Verbalized confidence obtains such an answer-level score by asking the model to emit a numerical probability with its final answer. Unlike confidence estimators based on token probabilities, hidden states, or repeated sampling (Luo et al., 2026; Srey et al., 2026; Wang & Stengel-Eskin, 2026), it can be produced in a single black-box generation (Xiong et al., 2024). Recent work trains this output with supervised fine-tuning (SFT), using targets derived from correctness or rollout statistics (Lin et al., 2022; Li et al., 2025), and with reinforcement learning (RL), by adding calibration terms to answer rewards (Xu et al., 2024; Damani et al., 2026; Ma et al., 2026; Yang et al., 2026).

We identify one primary failure and one recipe-dependent optimization effect. Under a fixed prompt, format, and decoding protocol, the base model’s sampled confidence is concentrated on a small set of values and only moderately stratifies correct and incorrect rollouts. The resulting confidence prior provides limited support for learning a graded correctness estimate. In filtering ablations, completed no-filtering DCPO and RLCR runs further concentrate confidence on a few values and endpoints, whereas the corresponding DAPO-filtered runs retain broader support. We therefore treat confidence concentration as a property of the sampling recipe, rather than as an inevitable consequence of GRPO.

We address the initialization problem with CalibSFT. CalibSFT constructs a target for each rollout by combining the success rate of its prompt group with the rollout’s verified correctness. Its selective segment loss retains confidence supervision on incorrect rollouts while restricting reasoning/answer supervision to correct rollouts, and target-balanced sampling exposes the model to a broader range of confidence values. In the subsequent RL stage, all compared methods use the same DAPO-style mixed-outcome filtering (Yu et al., 2025) as a fixed sampling recipe, allowing us to compare Base and CalibSFT initialization under controlled downstream optimization.

We make three contributions:

- We characterize a concentrated base-model confidence prior and document how confidence concentration depends on the RL sampling recipe.
- We introduce CalibSFT, which combines rollout-derived targets, target-balanced sampling, and correctness-aware segment supervision for confidence initialization.
- We evaluate whether the CalibSFT initialization improves accuracy, calibration, and confidence diversity after downstream RL across multiple confidence-aware objectives.

2 PRELIMINARIES

2.1 VERBALIZED CONFIDENCE FOR REASONING LLMs

In this work, we study confidence calibration for reasoning language models that verbalize a numerical confidence score (Tian et al., 2023; Xiong et al., 2024). Given a question x , the LLM is tasked with producing a completion $o = (r, a, c)$, where r is a reasoning trace and a is an answer. $c \in [0, 1]$ is a post-answer confidence value and represents the model’s estimated probability that a is correct.

Reinforcement Learning with Verifiable Rewards (RLVR). RLVR formulates the LLM as a policy π_θ parameterized by θ . Given a question x , the policy generates a rollout o , and a task verifier assigns it a binary correctness label $z = V(x, a) \in \{0, 1\}$. Typically, the policy is optimized using Group Relative Policy Optimization (GRPO) (Shao et al., 2024). In practice, GRPO samples G rollouts $\{o_i\}_{i=1}^G$ for each question, and the optimization objective can then be formulated as:

$$J_{\text{GRPO}}(\theta) = \mathbb{E} \left[\frac{1}{G} \sum_{i=1}^G \frac{1}{L_i} \sum_{t=1}^{L_i} \min \left(\rho_{i,t}(\theta) \hat{A}_i, \text{clip}(\rho_{i,t}(\theta), 1 - \epsilon_{\text{clip}}, 1 + \epsilon_{\text{clip}}) \hat{A}_i \right) \right]. \quad (1)$$

Here $o_i = (o_{i,1}, \dots, o_{i,L_i})$, and $\rho_{i,t}(\theta) = \pi_\theta(o_{i,t} | x, o_{i,<t}) / \pi_{\theta_{\text{old}}}(o_{i,t} | x, o_{i,<t})$ is the policy ratio; ϵ_{clip} is the clipping parameter. Given the reward $R_i = R(o_i)$ and a small constant $\epsilon > 0$ for numerical stability, the group-relative advantage is:

$$\hat{A}_i = \frac{R_i - \bar{R}}{\sigma_R + \epsilon}, \quad \bar{R} = \frac{1}{G} \sum_{j=1}^G R_j, \quad \sigma_R = \text{std}(\{R_j\}_{j=1}^G). \quad (2)$$

Confidence-Aware Rewards. For verbalized confidence training, a confidence-based scoring term is typically added to the reward. For example, RLCR (Damani et al., 2026) adds the negative Brier loss to the binary correctness reward:

$$R_i = z_i - (c_i - z_i)^2, \quad (3)$$

where c_i denotes the reported confidence and $z_i = V(x, a_i) \in \{0, 1\}$ denotes the correctness label of rollout i . This reward favors correct answers while assigning higher scores to confidence values that better match their correctness labels, thereby increasing their probability under GRPO.

2.2 CALIBRATION METRICS

A useful verbalized confidence score should be both calibrated and discriminative (Eusebi et al., 2026). To formalize calibration, let C and Z denote the confidence and correctness random variables, respectively. A perfectly calibrated score satisfies $\Pr(Z = 1 | C = p) = p$ for every $p \in [0, 1]$. To quantify the degree of miscalibration, **Expected Calibration Error (ECE)** (Guo et al., 2017) is defined as $\mathbb{E}_C[|\Pr(Z = 1 | C) - C|]$. In practice, after partitioning N examples into B confidence bins $\{\mathcal{I}_b\}_{b=1}^B$, we estimate ECE as

$$\text{ECE} = \sum_{b=1}^B \frac{|\mathcal{I}_b|}{N} |\text{acc}(\mathcal{I}_b) - \text{conf}(\mathcal{I}_b)|. \quad (4)$$

Here, $\text{acc}(\mathcal{I}_b)$ and $\text{conf}(\mathcal{I}_b)$ are the mean correctness and confidence in bin b , respectively. While ECE measures calibration at the bin level, the **Brier score** (Brier, 1950) measures the squared error between confidence and correctness for each instance,

$$\text{Brier} = \frac{1}{N} \sum_{j=1}^N (c_j - z_j)^2, \quad (5)$$

where lower values indicate better probabilistic predictions for both metrics. These metrics assess calibration, but neither directly evaluates discrimination as a ranking property. For example, assigning every response a constant confidence equal to the overall accuracy yields zero ECE while giving correct and incorrect responses identical scores (Eusebi et al., 2026). We therefore measure the discrimination using **AUROC**, which evaluates whether correct responses receive higher confidence than incorrect ones.




Figure 1: Confidence distributions for correct and incorrect responses from three base models on DeepScaleR-Train. Each distribution is normalized separately; annotations report overall Pass@1, Top-3 mass, and the proportion of confidence outputs below 0.5, pooling both groups. High confidence persists even when answers are incorrect.

3 MOTIVATION

Confidence-aware RL methods typically add a calibration reward computed from sampled responses (Xu et al., 2024; Damani et al., 2026; Ma et al., 2026; Yang et al., 2026). Since this reward acts on whatever confidence values the policy already samples, we first characterize the model’s confidence distribution before RL, then examine how RL changes it during training.

Confidence priors concentrate on a few high values. We refer to the confidence distribution before RL as the model’s *confidence prior*. To characterize this prior, we sample 10 rollouts per question from Qwen3-8B (Yang et al., 2025), Gemma4-E2B-Instruct (Abd et al., 2026), and OLMo3-7B-Instruct (Ettinger et al., 2025) on DeepScaleR-Train. Figure 1 presents their confidence distributions for correct and incorrect responses. All three models concentrate their confidence on a few high values, with fewer than 5% of outputs below 0.5. These base models also assign high confidence even to incorrect responses, suggesting that their verbalized confidence does not reliably reflect answer correctness. As a result, RL starts from a sub-optimal confidence prior that rarely produces low-confidence samples.

RL improves calibration but confidence concentration persists. To examine whether calibration-aware RL overcomes this concentration, we train Qwen3-8B with RLCR. As shown in Figure 2(a), RLCR reduces Brier loss over training, showing that the calibration reward can be optimized. However, Figure 2(b) shows that the model continues to sample a small set of confidence values for each question, with correspondingly low confidence token entropy throughout training (see Appendix B.2 for further training statistics).

To better understand why this concentration persists during training, we formulate confidence generation as action selection under a softmax policy (Garg et al., 2022). Since the LLM decodes each confidence value autoregressively, the induced distribution over the finite set of confidence values is itself a softmax distribution.

Proposition 1 (Policy gradients under confidence concentration). *Given a question x and a finite set of confidence actions $\{c_1, \dots, c_K\} \subseteq [0, 1]$, let $\ell = (\ell_1, \dots, \ell_K)$ denote the policy’s logits for these actions and $\pi_i = \Pr(C = c_i | x) = \exp(\ell_i) / \sum_{j=1}^K \exp(\ell_j)$. For $Z \in \{0, 1\}$ indicating the correctness of a rollout sampled for x , define the fixed expected reward $r_i = -\mathbb{E}[(c_i - Z)^2 | x]$. Let $J_x(\ell) = \sum_i \pi_i r_i$ and $\Delta_x = \max_i r_i - \min_i r_i$. Then*

$$\frac{\partial J_x}{\partial \ell_i} = \pi_i(r_i - J_x), \quad \left| \frac{\partial J_x}{\partial \ell_i} \right| \leq \pi_i \Delta_x. \quad (6)$$

If $\pi_k = 1 - \varepsilon$ for some $k \in \{1, \dots, K\}$, then

$$\|\nabla_{\ell} J_x\|_1 \leq 2\varepsilon \Delta_x. \quad (7)$$




Figure 2: Training dynamics of Qwen3-8B with RLCR. (a) Brier loss. (b) Distinct confidence values per question (eight rollouts; blue, left axis) and confidence token entropy (orange, right axis). Bold curves show centered five-step moving averages; faint curves show per-step measurements. Figure 3: Test-time calibration at step 160 on Math500 and MinervaMath (ID), using four rollouts per question. Bars show Pass@1 and mean confidence; task labels report Top-3 mass.

Proposition 1 shows that a rarely sampled confidence value has a small expected logit gradient even when choosing it would improve the expected reward. As the confidence distribution approaches a single choice, the entire logit gradient vanishes, including when that choice is suboptimal. This identifies how an initially concentrated confidence distribution can impede exploration of better-calibrated outputs, which is consistent with the persistently low distinct-value count and confidence token entropy observed in Figure 2(b). We provide the proof in Appendix B.3.

Persistent confidence concentration hinders downstream calibration. To examine the implications of persistent confidence concentration for downstream calibration, we evaluate the final RLCR checkpoint on two datasets of contrasting difficulty: Math500 and MinervaMath. As shown in Figure 3, the post-RL model underestimates its accuracy on Math500 (90.5% Pass@1 vs. 84.5% mean confidence) but is markedly overconfident on MinervaMath (36.7% Pass@1 vs. 74.8% mean confidence). Despite a 53.8-point gap in Pass@1 between the two tasks, mean confidence differs by only 9.7 points. The three most frequent values still account for 82.25% and 69.49% of outputs on Math500 and MinervaMath, respectively. The same small set of confidence values therefore appears to persist across tasks of different difficulty, failing to calibrate downstream tasks.

To characterize the cost of this miscalibration, we examine the expected Brier loss, which is minimized when confidence equals the conditional probability of correctness (Gneiting & Raftery, 2007; Damani et al., 2026).

Proposition 2 (Brier risk under limited confidence support). *Let $Z \in \{0, 1\}$ indicate answer correctness and let $C \in [0, 1]$ denote the predicted confidence. Given a question x , define $p_x = \Pr(Z = 1 \mid x)$. Assume that C and Z are conditionally independent given x , consistent with the limited discrimination between correct and incorrect responses observed in Figure 1. For any $\delta > 0$, let $m_\delta(x) = \Pr(|C - p_x| < \delta \mid x)$. Then*

$$\mathbb{E}[(C - Z)^2 \mid x] \geq p_x(1 - p_x) + \delta^2(1 - m_\delta(x)). \quad (8)$$

Proposition 2 shows that the lower bound on Brier loss increases as the probability that confidence falls near p_x decreases. Therefore, confidence concentration can hinder calibration when it leaves little probability mass near the probability of correctness. In Appendix B.4, we provide the detailed proof, extend the bound to a restricted confidence support, and show that limited support similarly raises the expected error of the closest of G sampled confidence values.

4 METHOD

In our previous analysis, we showed that confidence concentration persists during RL, since rarely sampled confidence values receive little gradient signal. This persistent concentration further hinders downstream calibration. To address this problem, our key idea is to supervise the confidence prior directly before RL, so that these values are learned rather than left for RL to explore. We propose

CalibSFT, a supervised fine-tuning stage for confidence initialization. For each question, CalibSFT samples multiple responses from the base model and assigns each a confidence target that combines the question’s success rate with the response’s own correctness. It learns confidence from both correct and incorrect responses, while using correct responses to supervise reasoning and answer generation.

4.1 PROPER-SCORING CONFIDENCE TARGETS

Binary correctness provides feedback on a particular answer, while repeated rollouts reveal how reliably the model solves the question. For each training prompt x_g , we sample K_{SFT} responses from the base model and use the task verifier to obtain correctness labels $z_{g,i} \in \{0, 1\}$. Following the group–instance target used in DCPO (Ma et al., 2026), we combine the success rate across these responses with each response’s correctness to construct SFT targets:

$$q_g = \frac{1}{K_{\text{SFT}}} \sum_{i=1}^{K_{\text{SFT}}} z_{g,i}, \quad c_{g,i}^* = \lambda q_g + (1 - \lambda) z_{g,i}. \quad (9)$$

The shared term q_g lets targets vary with the question’s difficulty for the base model, while $z_{g,i}$ distinguishes correct and incorrect responses within the group. The mixing weight $\lambda \in [0, 1]$ controls their relative contributions; we use $\lambda = 0.5$.

Combining these signals also reduces the variability of binary supervision. Both q_g and $z_{g,i}$ estimate the question’s success probability under the base model. We quantify how mixing them preserves this expectation and how the number of rollouts affects the target variance.

Proposition 3 (Mean and variance of confidence targets). *Given a prompt x_g and a sampling policy, let $Z_{g,1}, \dots, Z_{g,K}$ be conditionally independent $\text{Bernoulli}(p_g)$ correctness labels, where $p_g = \Pr(Z = 1 \mid x_g)$ and $K = K_{\text{SFT}}$. Define $Q_g = K^{-1} \sum_{j=1}^K Z_{g,j}$ and $T_{g,i} = \lambda Q_g + (1 - \lambda) Z_{g,i}$ for $\lambda \in [0, 1]$. Then*

$$\begin{aligned} \mathbb{E}[T_{g,i} \mid x_g] &= p_g, \\ \text{Var}(T_{g,i} \mid x_g) &= p_g(1 - p_g) \left[\frac{1}{K} + (1 - \lambda)^2 \left(1 - \frac{1}{K} \right) \right]. \end{aligned} \quad (10)$$

The confidence prediction that minimizes the expected squared loss given x_g is therefore

$$\arg \min_{a \in [0,1]} \mathbb{E}[(a - T_{g,i})^2 \mid x_g] = p_g. \quad (11)$$

Proposition 3 shows that mixing preserves the prompt-level optimum of Brier loss (Gneiting & Raftery, 2007). The target variance is no greater than that of an individual correctness label. Increasing the rollout count stabilizes the group success rate, while $\lambda < 1$ preserves the distinction between correct and incorrect responses. The proof and extension to general proper scoring rules are provided in Appendix C.

We replace each retained response’s confidence with $c_{g,i}^*$, preserving its reasoning and answer. To give each available target equal representation during SFT, we group responses by target value and sample the same number from each nonempty group. This gives infrequent targets the same training exposure as common ones. We provide filtering, formatting, and sampling details in Appendix D.1.

4.2 OUTCOME-CONDITIONAL SUPERVISION

The constructed data include incorrect responses paired with lower confidence targets. We retain these examples to teach the model to express uncertainty about its own errors. We use their reasoning and answers as context and supervise the confidence that follows. Correct responses receive supervision on the entire completion, reinforcing successful reasoning and answer generation alongside confidence prediction.

Formally, let $\tilde{o}$ denote a response with its confidence replaced by the constructed target, and let M and C denote its reasoning-and-answer and confidence token positions, respectively. We average the

token loss within each segment before combining the two supervision terms:

$$\begin{aligned} \mathcal{L}_{\text{CalibSFT}} = (1 - \beta) \mathbb{E}_{\text{correct}} & \left[-\frac{1}{|M|} \sum_{t \in M} \log \pi_{\theta}(\tilde{o}_t \mid x, \tilde{o}_{<t}) \right] \\ & + \beta \mathbb{E}_{\text{all}} \left[-\frac{1}{|C|} \sum_{t \in C} \log \pi_{\theta}(\tilde{o}_t \mid x, \tilde{o}_{<t}) \right]. \end{aligned} \quad (12)$$

The first expectation is over correct responses in the balanced SFT data, and the second is over all responses in that data. Separate segment averages prevent the longer reasoning sequences from dominating the loss by token count, and β controls the weight of confidence supervision. We use $\beta = 0.5$. The resulting checkpoint initializes confidence-aware RL, with each downstream method using its original reward and training configuration.

5 EXPERIMENTS

5.1 EXPERIMENTAL SETUP

Baselines. We compare methods spanning answer-only and calibration-aware training. **Base** denotes the checkpoint without task-specific training, while **RLVR** optimizes only answer correctness. We use confidence-aware RL baselines including **RLCR** (Damani et al., 2026), **ReDoubt** (Bani-Harouni et al., 2026), **CoCA** (Li et al., 2026), **DCPO** (Ma et al., 2026), and **C3RL** (Yang et al., 2026). To measure the effect of **CalibSFT**, we evaluate each confidence-aware RL method from both the Base and CalibSFT checkpoints. All methods use Qwen3-8B (Yang et al., 2025) as the default model and follow the **DAPO** training recipe (Yu et al., 2025).

Datasets. We use **DeepScaleR-Preview** (Luo et al., 2025) as the mathematical training dataset. We evaluate on a broad range of in-distribution (ID) and out-of-domain (OOD) datasets spanning different levels of difficulty. For ID mathematical reasoning, we use **DeepScaleR-Eval**, **MATH-500** (Hendrycks et al., 2021), **MinervaMath** (Lewkowycz et al., 2022), **OlympiadBench** (He et al., 2024), **GSM8K** (Cobbe et al., 2021), and **AIME 2024–2026** (Dekoninck et al., 2026). For general reasoning, we use **HotpotQA** (Yang et al., 2018), **MuSiQue** (Trivedi et al., 2022), **TriviaQA** (Joshi et al., 2017), **NQOpen** (Lee et al., 2019), **PopQA** (Mallen et al., 2023), **WebQuestions** (Berant et al., 2013), **DROP** (Dua et al., 2019), and **LiveBenchReasoning** (White et al., 2024).

Evaluation metrics. We evaluate model performance using complementary metrics that capture answer correctness, confidence discrimination, and absolute calibration. We report **Pass@1** for answer correctness, **AUROC** for ranking correct responses above incorrect responses, and **Brier score** (Brier, 1950) and **ECE** (Guo et al., 2017) for probabilistic calibration.

Implementation details. During evaluation, we use temperature 0.6 and sample 32 rollouts per problem for the AIME benchmarks and 4 rollouts per problem for all other benchmarks. All experiments are conducted on a single node with eight NVIDIA A100 GPUs. We report our hyperparameter setting and training details in Appendix D.

5.2 OVERALL RESULTS

CalibSFT improves calibration across diverse confidence-aware RL methods. As shown in Table 1, initializing RL with CalibSFT reduces Brier score and ECE on mathematical reasoning benchmarks for all three completed paired comparisons: RLCR, CoCA, and DCPO. AUROC also improves for all three methods. We provide detailed per-dataset results in Table 8 (Appendix E.1). For example, CalibSFT reduces DCPO’s ECE from 44.59 to 20.25 while increasing its AUROC from 64.19 to 74.73 on MinervaMath. CalibSFT achieves these gains while largely preserving answer accuracy, with an average Pass@1 gain of 1.32 points across these three RL methods. These results support its use as a plug-and-play initialization for better confidence-aware RL.

CalibSFT generalizes to general reasoning tasks beyond the training domain. To evaluate the generalization of CalibSFT beyond its mathematical domain, we assess the resulting models on

Table 1: Average performance across 8 ID and 8 OOD datasets. Higher is better for Pass@1 and AUROC; lower is better for Brier and ECE. Each entry averages the datasets within the corresponding split. Base, RLVR, and standalone CalibSFT use a matched 6,144-token evaluation budget.

Method	In-distribution				Out-of-distribution			
	Pass@1 $\uparrow$	AUROC $\uparrow$	Brier $\downarrow$	ECE $\downarrow$	Pass@1 $\uparrow$	AUROC $\uparrow$	Brier $\downarrow$	ECE $\downarrow$
Base	57.39	71.47	34.97	36.26	45.36	65.71	45.15	46.94
RLVR	61.60	73.16	33.48	34.08	43.88	64.19	50.29	51.52
CalibSFT	56.20	79.25	18.01	14.82	45.91	61.38	28.25	24.16
RLCR	62.71	82.78	15.60	16.84	43.96	69.38	31.14	31.17
+ CalibSFT	64.70	86.88	13.48	10.84	45.71	71.31	23.24	18.20
ReDoubt	63.08	81.70	15.94	18.85	43.71	69.88	25.07	23.98
+ CalibSFT	64.99	88.33	13.00	16.20	45.55	73.93	22.42	18.89
CoCA	61.11	80.59	17.25	21.75	43.68	66.79	26.67	24.12
+ CalibSFT	62.65	86.31	16.40	20.21	45.55	71.21	22.95	18.89
DCPO	61.98	78.16	16.38	14.68	43.99	68.13	32.85	32.41
+ CalibSFT	62.40	85.16	13.29	12.20	44.82	72.75	21.73	17.31
C3RL	63.22	67.90	27.80	30.45	43.44	62.19	45.54	48.00
+ CalibSFT	63.32	80.88	15.65	15.18	44.60	63.03	23.89	17.38

Table 2: Correlation between accuracy and mean verbalized confidence across ID and OOD datasets. Base and +Ours denote RL without and with CalibSFT initialization, respectively; the higher value is bold. Each coefficient uses eight equally weighted datasets within its split.

Dataset	Metric	RLCR		ReDoubt		CoCA		DCPO		C3RL	
		Base	+Ours	Base	+Ours	Base	+Ours	Base	+Ours	Base	+Ours
ID	Pearson r	0.768	0.963	0.829	0.976	0.750	0.887	0.669	0.996	0.635	0.896
	Spearman ρ	0.619	0.952	0.762	0.976	0.643	0.976	0.619	0.976	0.595	0.833
OOD	Pearson r	0.505	0.729	0.667	0.738	0.447	0.605	0.467	0.611	0.681	0.848
	Spearman ρ	0.548	0.619	0.762	0.643	0.571	0.619	0.548	0.643	0.714	0.929

eight general reasoning benchmarks. Table 1 shows lower average Brier score and ECE for all three completed paired comparisons on these benchmarks. For example, CalibSFT reduces DCPO’s ECE from 20.53 to 8.55 on MuSiQue. We provide detailed per-dataset results in Table 9 (Appendix E.1).

CalibSFT improves confidence–accuracy correlation across tasks of varying difficulty. To assess whether average confidence tracks accuracy across tasks, we compare each dataset’s average confidence with its accuracy. For each RL method with and without CalibSFT initialization, we compute Pearson and Spearman correlations separately over the eight ID and eight OOD datasets. Table 2 shows that CalibSFT improves Spearman correlation for all three completed RL comparisons on mathematical and general reasoning benchmarks. For example, DCPO’s Spearman correlation increases from 0.619 to 0.976 on mathematical tasks and from 0.548 to 0.643 on general reasoning tasks. Pearson correlation also improves for DCPO and for RLCR on ID tasks, while remaining essentially unchanged for RLCR on OOD tasks. These results suggest that CalibSFT helps downstream RL models assign higher average confidence to tasks on which they perform better.

CalibSFT mitigates confidence concentration after downstream RL. Prior work shows that confidence estimates can remain concentrated on a few values even after calibration-aware RL training (Wang et al., 2026; Yang et al., 2026). To examine how CalibSFT initialization influences confidence distribution, we compare RLCR, CoCA, and DCPO on MinervaMath. As shown in Figure 4, models trained without CalibSFT tend to concentrate their confidence on a few values and show higher average confidence than their accuracy. CalibSFT reduces the dominant peaks and brings average confidence closer to accuracy, alleviating both concentration and overconfidence. For




Figure 4: Confidence distributions on MinervaMath for RLCR, CoCA, and DCPO, with four rollouts per question at temperature 0.6. CalibSFT reduces concentration at dominant confidence values and lowers ECE for all three methods.




Figure 5: Accuracy (left) and ECE (right) on AIME’26 across token budgets for RLCR with and without CalibSFT initialization. CalibSFT yields substantially lower ECE across all budgets while matching accuracy at medium budgets and improving it at larger budgets.

Table 3: The calibration comparison between pre-SFT constructed targets and post-SFT predictions. CalibSFT achieves higher AUROC and lower Brier score across ID mathematical benchmarks.

example, DCPO’s average confidence decreases from 79.32% to 63.69%, while accuracy increases from 34.93% to 36.58%. ECE decreases for all three methods, from 44.59% to 28.35% for DCPO. We provide additional confidence distribution metrics in Appendix E.7.

The calibration gains of CalibSFT persist across reasoning budgets. To assess whether CalibSFT improves calibration as reasoning progresses rather than the final answer alone, we compare RLCR with and without CalibSFT initialization on AIME 2026 under varying token budgets. As shown in Figure 5, CalibSFT yields substantially lower ECE across all token budgets. It matches the RLCR baseline in accuracy around 5,000 tokens and becomes more accurate at larger budgets. For example, CalibSFT reduces ECE from 22.71 to 10.97 at a 10,000-token budget, while accuracy increases from 47.50% to 52.71%. These results suggest that CalibSFT’s calibration benefits persist as more computation is allocated to reasoning.

5.3 ABLATION STUDIES

Confidence target construction. The construction of confidence targets in CalibSFT is inspired by proper scoring rules, which serve as a foundation in calibration. To verify the effectiveness of this target design, we compare CalibSFT with targets based on group-level correctness (CSFT) (Jang et al., 2025), token probabilities (UaIT) (Liu et al., 2024), semantic consistency (SaySelf) (Xu et al., 2024), and rubric-based LLM judgments (LoVeC) (Zhang et al., 2026), using a shared SFT training set the same optimization objective. We provide details of the subset and target construction in Appendix E.2. We report the results of the target construction in Table 3, which shows the AUROC and Brier score of the constructed targets before SFT and average performance on ID mathematical benchmarks after SFT. CalibSFT yields lower Brier score and higher AUROC than the four alternatives. These results

Table 5: Supervision settings and average performance on eight ID benchmarks. A check mark indicates that a segment contributes to the training loss; segment balancing averages the reasoning/answer and confidence losses separately. All variants use the same statically balanced training set and train for 180 steps.

Variant	Reasoning/answer		Confidence		Segment balanced	ID performance			
	Correct	Incorrect	Correct	Incorrect		Pass@1 $\uparrow$	AUROC $\uparrow$	Brier $\downarrow$	ECE $\downarrow$
Confidence-only	–	–	✓	✓	–	43.56	74.62	20.97	17.43
Full-CE	✓	✓	✓	✓	–	65.89	61.66	29.29	27.65
Full-Balanced	✓	✓	✓	✓	✓	65.46	80.67	17.04	13.87
Correct-only	✓	–	✓	–	✓	66.59	81.58	16.95	14.28
CalibSFT	✓	–	✓	✓	✓	65.79	79.30	17.89	14.36

support the effectiveness of our confidence target construction, which accounts for both calibration and discrimination.

Target mixing weight. We construct confidence targets as $c_{g,i}^* = \lambda q_g + (1 - \lambda)z_{g,i}$, where q_g is the group success rate and $z_{g,i}$ is the correctness of an individual response. To evaluate the effect of the mixing weight, we compare $\lambda \in \{0, 0.25, 0.5, 0.75, 1\}$ using the same subset as above and keeping the SFT settings fixed. At $\lambda = 0$, targets are restricted to 0 and 1; this setting yields the highest Brier score and ECE in Table 4. At $\lambda = 1$, correct and incorrect responses to the same question receive identical targets, providing no within-question discrimination signal; post-SFT AUROC is also lower than at $\lambda = 0.5$. Equal weighting achieves the highest AUROC and lowest Brier score and ECE among the tested settings, so we use $\lambda = 0.5$ as the default.

Table 4: Average calibration performance on ID benchmarks across mixing weights. Equal weighting ($\lambda = 0.5$) performs best on all three metrics.

λ	AUROC $\uparrow$	Brier $\downarrow$	ECE $\downarrow$
0	68.10	32.49	32.49
0.25	79.02	18.83	14.21
0.5	79.75	17.81	13.03
0.75	65.93	19.99	13.16
1	73.29	19.90	15.04

Supervision design. [TODO:] Compare confidence-only SFT, full-completion CE, segment-balanced full-completion CE, CalibSFT, and CalibSFT without target-balanced sampling. Hold the training examples, confidence targets, and optimization settings fixed across the supervision variants. Report which completion segments contribute to each objective, followed by their ID/OOD Pass@1 and calibration results. Use this ablation to isolate the effects of retaining answer supervision on correct rollouts, masking incorrect reasoning and answers, supervising confidence on all rollouts, and balancing confidence targets.

6 DISCUSSION

6.1 VERBALIZED CONFIDENCE AND ALTERNATIVE CONFIDENCE SIGNALS

Verbalized confidence provides a directly interpretable uncertainty signal. [TODO:] Compare post-answer verbalized confidence with token- or sequence-level probability, self-consistency, semantic consistency, and a hidden-state probe. Use the same generated answers whenever the signal permits, and report AUROC, Brier score, ECE, inference cost, and model-access requirements. Discuss which gains arise from CalibSFT and which arise from the choice of confidence signal.

6.2 CONFIDENCE FORMATS AND RESOLUTION

CalibSFT can be instantiated with different confidence formats. [TODO:] Compare two-decimal probabilities, discrete confidence levels, and semantic labels such as *low*, *medium*, and *high*. Map all formats to $[0, 1]$ for evaluation and report Pass@1, AUROC, Brier score, and ECE. Include

ConfTuner as the representative tokenized-confidence method when its output format and evaluation mapping are matched.

The rollout count controls the resolution of empirical targets. [TODO:] Fix $\lambda = 0.5$ and compare target construction with $K_{\text{SFT}} \in \{10, 20, 100\}$. Report the resulting target coverage, target estimation cost, and final calibration. Use this comparison to separate the effect of a finer confidence grid from the effect of additional rollout samples.

7 DOWNSTREAM APPLICATIONS

7.1 SELECTIVE ANSWERING AND RISK CONTROL

[TODO:] Rank responses by confidence and report risk–coverage curves, AURC, and accuracy at matched coverage. Compare Base, confidence-aware RL, and CalibSFT-initialized RL to test whether the calibration improvements translate into more reliable abstention decisions.

7.2 CONFIDENCE-GUIDED EFFICIENT REASONING

[TODO:] Use confidence to allocate additional generation budget or route uncertain questions to a stronger inference procedure. Compare against fixed-budget and random-routing baselines at matched token cost, reporting accuracy, calibration, and total generated tokens.

8 CONCLUSION

[TODO:] Summarize the concentrated and misaligned confidence prior, the rollout-derived and target-balanced CalibSFT initialization, and its effect on downstream confidence-aware RL. Close with the main ID/OOD evidence and the remaining limitations of rollout-derived confidence targets.

AI USE STATEMENT

(This section is **required** and does not count toward the page limit.)

In this work, we used generative AI tools for [tasks with required disclosure]. We have not used generative AI tools for [other tasks with required disclosure], and [the rest of the required disclosure tasks] are not applicable to this work. Additionally, we used generative AI tools for [tasks with recommended, disclosure]. We have reviewed all AI-assisted work. [Elaborate. For example, “we checked LLM-generated research ideas for potential plagiarism through a manual literature survey”, “LLM-generated code was verified and tested for correctness by 2 authors”, etc.]. We take responsibility for the final content of this work, including text, claims or artifacts produced with the aid of generative AI.

See the ICLR 2027 AI Policy for Authors for more details. This statement should not be more than 1 page.

ETHICS STATEMENT

(This section is **recommended** and does not count toward the page limit.)

If authors feel that their paper submission raises questions regarding the Code of Ethics, they are encouraged to include a paragraph of Ethics Statement (at the end of the main text before references) to address potential concerns where appropriate. Topics include, but are not limited to, studies that involve human subjects, practices to data set releases, potentially harmful insights, methodologies and applications, potential conflicts of interest and sponsorship, discrimination/bias/fairness concerns, privacy and security issues, legal compliance, and research integrity issues (e.g., IRB, documentation, research ethics). This statement should not be more than 1 page.

REPRODUCIBILITY STATEMENT

(This section is **recommended** and does not count toward the page limit.)

It is important that the work published in ICLR is reproducible. Authors are strongly encouraged to include a paragraph-long Reproducibility Statement at the end of the main text (before references) to discuss the efforts that have been made to ensure reproducibility. This paragraph should not itself describe details needed for reproducing the results, but rather reference the parts of the main paper, appendix, and supplemental materials that will help with reproducibility. For example, for novel models or algorithms, a link to an anonymous downloadable source code can be submitted as supplementary materials; for theoretical results, clear explanations of any assumptions and a complete proof of the claims can be included in the appendix; for any datasets used in the experiments, a complete description of the data processing steps can be provided in the supplementary materials. Each of the above are examples of things that can be referenced in the reproducibility statement.

AUTHOR CONTRIBUTIONS

If you'd like to, you may include a section for author contributions as is done in many journals. This is optional and at the discretion of the authors.

ACKNOWLEDGMENTS

Use unnumbered third level headings for the acknowledgments. All acknowledgments, including those to funding agencies, go at the end of the paper.

REFERENCES

- Sherif El Abd, Vaibhav Aggarwal, Robin Algayres, Alek Andreev, Olivier Bachem, Ian Ballantyne, Cormac Brick, Victor Cărbune, Michelle Casbon, et al. Gemma 4 technical report. *arXiv preprint arXiv:2607.02770*, 2026.
- Amos Azaria and Tom Mitchell. The internal state of an llm knows when it's lying. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pp. 967–976, 2023.
- David Bani-Harouni, Chantal Pellegrini, Paul Stangel, Ege Özsoy, Kamilia Zaripova, Nassir Navab, and Matthias Keicher. Rewarding doubt: A reinforcement learning approach to calibrated confidence expression of large language models. In *The Fourteenth International Conference on Learning Representations*, 2026.
- Jonathan Berant, Andrew Chou, Roy Frostig, and Percy Liang. Semantic parsing on Freebase from question-answer pairs. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*, pp. 1533–1544, 2013.
- Glenn W. Brier. Verification of forecasts expressed in terms of probability. *Monthly Weather Review*, 78(1):1–3, 1950.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Mehul Damani, Isha Puri, Stewart Slocum, Idan Shenfeld, Leshem Choshen, Yoon Kim, and Jacob Andreas. Beyond binary rewards: Training LMs to reason about their uncertainty. In *International Conference on Learning Representations*, 2026.
- Jasper Dekoninck, Nikola Jovanović, Tim Gehrunger, Kári Rognvaldsson, Ivo Petrov, Chenhao Sun, and Martin Vechev. Beyond benchmarks: Matharena as an evaluation platform for mathematics with LLMs. In *3rd AI for Math Workshop: Toward Self-Evolving Scientific Agents, ICML Workshop AI4Math*, 2026.
- Dheeru Dua, Yizhong Wang, Pradeep Dasigi, Gabriel Stanovsky, Sameer Singh, and Matt Gardner. DROP: A reading comprehension benchmark requiring discrete reasoning over paragraphs. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pp. 2368–2378, 2019.

- Allyson Ettinger, Amanda Bertsch, Bailey Kuehl, et al. Olmo 3. *arXiv preprint arXiv:2512.13961*, 2025. URL <https://arxiv.org/abs/2512.13961>.
- Aliai Eusebi, Alexander Herzog, Xiaoyu Liang, Marie Vasek, Enrico Mariconti, and Lorenzo Cavallaro. Reading calibrated uncertainty from language model trajectories. *arXiv preprint arXiv:2605.22864*, 2026.
- Conor Finlay, Joshua Kurien, Saurabh Dash, Marzieh Fadaee, and Beyza Ermis. Caliber: Calibrating confidence before and after reasoning in language models. *arXiv preprint arXiv:2606.24281*, 2026.
- Shivam Garg, Samuele Tosatto, Yangchen Pan, Martha White, and A. Rupam Mahmood. An alternate policy gradient estimator for softmax policies, 2022.
- Jiahui Geng, Fengyu Cai, Yuxia Wang, Heinz Koepl, Preslav Nakov, and Iryna Gurevych. A survey of confidence estimation and calibration in large language models. In *Proceedings of the 2024 conference of the north American chapter of the association for computational linguistics: human language technologies (volume 1: long papers)*, pp. 6577–6595, 2024.
- Tilmann Gneiting and Adrian E. Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007. doi: 10.1198/016214506000001437.
- Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning*, pp. 1321–1330, 2017.
- Chaoqun He, Renjie Luo, Yuzhuo Bai, Shengding Hu, Zhen Leng Thai, Junhao Shen, Jinyi Hu, Xu Han, Yujie Huang, Yuxiang Zhang, Jie Liu, Lei Qi, Zhiyuan Liu, and Maosong Sun. Olympiad-bench: A challenging benchmark for promoting AGI with olympiad-level bilingual multimodal scientific problems. *arXiv preprint arXiv:2402.14008*, 2024.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In *Advances in Neural Information Processing Systems*, 2021.
- Chaeyun Jang, Moonseok Choi, Yegon Kim, Hyungi Lee, and Juho Lee. Verbalized confidence triggers self-verification: Emergent behavior without explicit reasoning supervision. *arXiv preprint arXiv:2506.03723*, 2025. URL <https://arxiv.org/abs/2506.03723>.
- Mandar Joshi, Eunsol Choi, Daniel S. Weld, and Luke Zettlemoyer. TriviaQA: A large scale distantly supervised challenge dataset for reading comprehension. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics*, pp. 1601–1611, 2017.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.
- Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. Semantic uncertainty: Linguistic invariances for uncertainty estimation in natural language generation. *arXiv preprint arXiv:2302.09664*, 2023.
- Kenton Lee, Ming-Wei Chang, and Kristina Toutanova. Latent retrieval for weakly supervised open domain question answering. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pp. 6086–6096, 2019. doi: 10.18653/v1/P19-1612.
- Aitor Lewkowycz, Anders Andreassen, David Dohan, Ethan Dyer, Henryk Michalewski, Vinay Ramasesh, Ambrose Slone, Cem Anil, Imanol Schlag, Theo Gutman-Solo, Yuhuai Wu, Behnam Neyshabur, Guy Gur-Ari, and Vedant Misra. Solving quantitative reasoning problems with language models. In *Advances in Neural Information Processing Systems*, 2022.
- Changcheng Li, Jiancan Wu, Hengheng Zhang, Zhengsu Chen, Guo An, Junxiang Qiu, Xiang Wang, and Qi Tian. Confidence before answering: A paradigm shift for efficient llm uncertainty estimation. *arXiv preprint arXiv:2603.05881*, 2026.

- Yibo Li, Miao Xiong, Jiaying Wu, and Bryan Hooi. ConfTuner: Training large language models to express their confidence verbally. *arXiv preprint arXiv:2508.18847*, 2025.
- Stephanie Lin, Jacob Hilton, and Owain Evans. Teaching models to express their uncertainty in words. *Transactions on Machine Learning Research*, 2022.
- Shudong Liu, Zhaocong Li, Xuebo Liu, Runzhe Zhan, Derek F. Wong, Lidia S. Chao, and Min Zhang. Can LLMs learn uncertainty on their own? expressing uncertainty effectively in a self-training manner. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 21635–21645, 2024. doi: 10.18653/v1/2024.emnlp-main.1205.
- Xiaoou Liu, Tiejun Chen, Longchao Da, Chacha Chen, Zhen Lin, and Hua Wei. Uncertainty quantification and confidence calibration in large language models: A survey. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V. 2*, pp. 6107–6117, 2025.
- Beier Luo, Shuoyuan Wang, Sharon Li, and Hongxin Wei. Your pre-trained llm is secretly an unsupervised confidence calibrator. *Advances in Neural Information Processing Systems*, 38: 163514–163543, 2026.
- Michael Luo, Sijun Tan, Justin Wong, Xiaoxiang Shi, William Y. Tang, Manan Roongta, Colin Cai, Jeffrey Luo, Li Erran Li, Raluca Ada Popa, and Ion Stoica. Deepscaler: Surpassing o1-preview with a 1.5b model by scaling rl, 2025. Notion Blog.
- Zhengzhao Ma, Xueru Wen, Boxi Cao, Yaojie Lu, Hongyu Lin, Jinglin Yang, Min He, Xianpei Han, and Le Sun. Decoupling reasoning and confidence: Resurrecting calibration in reinforcement learning from verifiable rewards. In *Forty-third International Conference on Machine Learning*, 2026.
- Alex Mallen, Akari Asai, Victor Zhong, Rajarshi Das, Daniel Khashabi, and Hannaneh Hajishirzi. When not to trust language models: Investigating effectiveness of parametric and non-parametric memories. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 9802–9822, 2023. doi: 10.18653/v1/2023.acl-long.546.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.
- Ponhvoan Srey, Xiaobao Wu, Cong-Duy Nguyen, Quang Minh Nguyen, Duc Anh Vu, and Anh Tuan Luu. From signals to transfer: A factorised study of probe-based uncertainty estimation in large language models. *arXiv preprint arXiv:2606.27679*, 2026.
- Katherine Tian, Eric Mitchell, Allan Zhou, Archit Sharma, Rafael Rafailov, Huaxiu Yao, Chelsea Finn, and Christopher D. Manning. Just ask for calibration: Strategies for eliciting calibrated confidence scores from language models fine-tuned with human feedback. *arXiv preprint arXiv:2305.14975*, 2023.
- Harsh Trivedi, Niranjan Balasubramanian, Tushar Khot, and Ashish Sabharwal. MuSiQue: Multihop questions via single-hop question composition. *Transactions of the Association for Computational Linguistics*, 10:539–554, 2022.
- Liaoyaqi Wang, Chunsheng Zuo, William Jurayj, Benjamin Van Durme, and Anqi Liu. Process supervision of confidence margin for calibrated LLM reasoning. In *Conference on Language Modeling*, 2026.
- Victor Wang and Elias Stengel-Eskin. Calibrating verbalized confidence with self-generated distractors. In *International Conference on Learning Representations*, volume 2026, pp. 121268–121297, 2026.
- Colin White, Samuel Dooley, Manley Roberts, Arka Pal, Ben Feuer, Siddhartha Jain, Ravid Shwartz-Ziv, Neel Jain, Khalid Saifullah, Siddhartha Naidu, Chinmay Hegde, Yann LeCun, Tom Goldstein, Willie Neiswanger, and Micah Goldblum. LiveBench: A challenging, contamination-free LLM benchmark. *arXiv preprint arXiv:2406.19314*, 2024. URL <https://arxiv.org/abs/2406.19314>.

- Jiayun Wu, Jiashuo Liu, Zhiyuan Zeng, Tianyang Zhan, Tianle Cai, and Wenhao Huang. Mitigating LLM hallucination via behaviorally calibrated reinforcement learning. In *Agentic AI in the Wild: From Hallucinations to Reliable Autonomy*, 2026.
- Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. Can LLMs express their uncertainty? an empirical evaluation of confidence elicitation in LLMs. In *International Conference on Learning Representations*, 2024.
- Tianyang Xu, Shujin Wu, Shizhe Diao, Xiaozhe Liu, Xingyao Wang, Yangyi Chen, and Jing Gao. SaySelf: Teaching LLMs to express confidence with self-reflective rationales. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 5985–5998, 2024.
- An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- Xuqing Yang, Yi Yuan, Shanzhe Lei, and Xuhong Wang. Scaling with confidence: Calibrating confidence of llms for adaptive test time scaling. *arXiv preprint arXiv:2607.01612*, 2026.
- Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William W. Cohen, Ruslan Salakhutdinov, and Christopher D. Manning. HotpotQA: A dataset for diverse, explainable multi-hop question answering. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pp. 2369–2380, 2018.
- Qiyang Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian Fan, Gaohong Liu, Lingjun Liu, et al. DAPO: An open-source LLM reinforcement learning system at scale. *arXiv preprint arXiv:2503.14476*, 2025.
- Caiqi Zhang, Xiaochen Zhu, Chengzu Li, Nigel Collier, and Andreas Vlachos. LoVeC: Reinforcement learning for better verbalized confidence in long-form generation. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 33336–33363. Association for Computational Linguistics, 2026. doi: 10.18653/v1/2026.acl-long.1539.

A RELATED WORK

A.1 CONFIDENCE ESTIMATION IN LLMs

LLM confidence can be estimated from token probabilities, hidden-state probes, agreement across sampled answers, or explicit self-reports (Kadavath et al., 2022; Azaria & Mitchell, 2023; Kuhn et al., 2023; Xiong et al., 2024). Verbalized-confidence methods ask the model to emit a numerical or linguistic confidence score, providing a single-generation estimate without requiring access to internal model representations (Tian et al., 2023; Xiong et al., 2024). Our work studies this post-answer verbalized setting.

A.2 CONFIDENCE-AWARE POST-TRAINING

Supervised approaches train confidence outputs using correctness labels, empirical rollout frequencies, or confidence-token objectives (Lin et al., 2022; Xu et al., 2024; Li et al., 2025). Calibration-aware RL instead incorporates proper scoring rewards or modifies the confidence target and credit assignment. Representative methods include RLCR, Rewarding Doubt, CoCA, DCPO, and CALIBER (Damani et al., 2026; Bani-Harouni et al., 2026; Li et al., 2026; Ma et al., 2026; Finlay et al., 2026). CalibSFT focuses on the supervised initialization supplied to these downstream RL objectives.

B ADDITIONAL ANALYSIS OF THE MOTIVATION

B.1 COMPLETE CONFIDENCE DISTRIBUTIONS

For the initial confidence analysis in Section 3, each model generates ten responses to each of the same 36,915 DeepScaleR-Train questions, with temperature 1.0, top- p 0.95, and an 8,192-token

Table 6: Overall accuracy and confidence concentration before RL. All entries are percentages; confidence statistics pool correct and incorrect responses.

Model	Pass@1	Top-3 mass	Confidence < 0.5
Qwen3-8B	62.01	93.94	0.61
Gemma4-E2B	54.99	91.06	1.64
OLMo3-7B	48.72	65.30	4.90

budget. We analyze the original confidence outputs before target construction. Valid-confidence rates are 80.57%, 99.81%, and 79.49% for Qwen3-8B, Gemma4-E2B-Instruct, and OLMo3-7B-Instruct, respectively. Distribution statistics and correctness-conditional confidence means use valid outputs. Figure 1 separates correct and incorrect responses, normalizing each group’s density to integrate to one. Confidence values are rounded to two decimal places before Gaussian smoothing with bandwidth 0.04 and reflection at zero and one. Table 6 reports the pooled statistics from the rounded, unsmoothed values; Pass@1 uses all rollouts.

B.2 CONFIDENCE SAMPLING DURING RL

Table 7 reports Base→RLCR statistics averaged over training steps. Confidence token entropy is the mean next-token entropy at confidence-value token positions. Top-3 share is the fraction of a step’s sampled confidence values assigned to its three most frequent values. Distinct confidence counts are computed within each question’s group of eight responses and then averaged. Statistics use the groups retained by the training filter; the sampled questions and responses can change during training. The logged confidence values include zero placeholders for missing outputs. Excluding all zero values, the final 40 steps still have a mean top-3 share of 45.16%.

Table 7: Brier loss and confidence sampling during Base→RLCR training. Later training lowers Brier loss while sampled confidence values remain concentrated.

Steps	Brier	Confidence token entropy	Top-3 share (%)	Distinct values per group
1–10	0.266	0.325	81.00	2.95
11–40	0.293	0.244	76.51	3.26
41–80	0.310	0.196	74.15	3.26
81–120	0.230	0.322	46.07	4.58
121–160	0.225	0.308	44.04	4.59

B.3 PROOF OF PROPOSITION 1

Proof. We fix x and the expected rewards r_i , and differentiate with respect to the confidence logits. The softmax derivative is $\partial\pi_j/\partial\ell_i = \pi_j(\mathbf{1}\{j = i\} - \pi_i)$. It follows that

$$\frac{\partial J_x}{\partial\ell_i} = \sum_j r_j \pi_j (\mathbf{1}\{j = i\} - \pi_i) \quad (13)$$

$$= \pi_i (r_i - J_x). \quad (14)$$

Since J_x is a convex combination of the r_i values, $|r_i - J_x| \leq \Delta_x$, which gives the coordinate-wise bound. If $\pi_k = 1 - \varepsilon$, then

$$|r_k - J_x| = \left| \sum_{i \neq k} \pi_i (r_k - r_i) \right| \leq \varepsilon \Delta_x. \quad (15)$$

Consequently,

$$\|\nabla_{\ell} J_x\|_1 = \pi_k |r_k - J_x| + \sum_{i \neq k} \pi_i |r_i - J_x| \quad (16)$$

$$\leq (1 - \varepsilon) \varepsilon \Delta_x + \varepsilon \Delta_x \quad (17)$$

$$\leq 2\varepsilon \Delta_x. \quad (18)$$

□

The gradient is the expectation of the likelihood-ratio estimator $R \nabla_{\ell} \log \pi_I$, where $I \sim \pi$ and $R = -(c_I - Z)^2$, with Z drawn from its fixed conditional distribution given x , independently of I . This matches RLCR's reward $Z - (c_I - Z)^2$ (Preliminaries) up to the term Z , which does not depend on c_I . Assuming $C \perp Z \mid x$ as in Proposition 2, this term contributes an ℓ -independent additive constant to J_x that does not affect this gradient. A baseline depending only on x leaves this expectation unchanged.

B.4 PROOF OF PROPOSITION 2

Proof. Conditional independence gives $\mathbb{E}[Z - p_x \mid x, C] = 0$. Expanding the square therefore yields

$$\mathbb{E}[(C - Z)^2 \mid x] = \mathbb{E}[(Z - p_x)^2 \mid x] + \mathbb{E}[(C - p_x)^2 \mid x] \quad (19)$$

$$= p_x(1 - p_x) + \mathbb{E}[(C - p_x)^2 \mid x]. \quad (20)$$

For any $\delta > 0$, the second term satisfies

$$\mathbb{E}[(C - Z)^2 \mid x] - p_x(1 - p_x) = \mathbb{E}[(C - p_x)^2 \mid x] \quad (21)$$

$$\geq \delta^2 \Pr(|C - p_x| \geq \delta \mid x) \quad (22)$$

$$= \delta^2(1 - m_{\delta}(x)). \quad (23)$$

Thus, assigning little probability mass near p_x incurs excess Brier loss even when confidence values in that neighborhood remain possible. □

Restricted confidence support. As a special case, suppose $\Pr(C \in S_x \mid x) = 1$ for a nonempty set $S_x \subseteq [0, 1]$, and define $d_x = \text{dist}(p_x, S_x) = \inf_{c \in S_x} |c - p_x|$. If $d_x > 0$, then $m_{d_x}(x) = 0$, and Proposition 2 with $\delta = d_x$ gives

$$\mathbb{E}[(C - Z)^2 \mid x] \geq p_x(1 - p_x) + \text{dist}(p_x, S_x)^2. \quad (24)$$

For $d_x = 0$, the same bound follows directly from the decomposition above. If S_x is closed, a predictor that always outputs a nearest point in S_x attains the bound. Confidence values are effectively drawn from a small discrete set throughout this paper (Figure 1), so $\text{dist}(p_x, S_x)$ measures how far the nearest available value falls from p_x , translating the observed concentration directly into excess Brier risk.

Finite-sampling limitation. Fix the policy and a question x , and draw G independent confidence samples $C_1, \dots, C_G$ from its conditional distribution. Each confidence sample lies outside $\{c : |c - p_x| < \delta\}$ with probability $1 - m_{\delta}(x)$. Independence gives the probability $(1 - m_{\delta}(x))^G$ that all G samples lie outside. On that event, even the sample closest to p_x has squared error at least δ^2 , so

$$\mathbb{E} \left[\min_{i \leq G} (C_i - p_x)^2 \mid x \right] \geq \delta^2 (1 - m_{\delta}(x))^G. \quad (25)$$

This statement fixes the policy and the question x ; it does not assume that the confidence support remains unchanged during RL.

B.5 ADDITIONAL MODELS AND SAMPLING SETTINGS

[TODO:] Add the motivation analysis for additional model families and sampling settings once those evaluation runs are complete.

C PROPERTIES OF THE MIXED CONFIDENCE TARGET

We analyze the unrounded targets under the original sampling distribution, before response filtering and target balancing. The random variables Q_g and $T_{g,i}$ correspond to the observed q_g and $c_{g,i}^*$ in Equation 9.

Proof of Proposition 3. Fix x_g and write $v_g = p_g(1 - p_g)$. Conditional independence gives

$$\mathbb{E}[Q_g | x_g] = \mathbb{E}[Z_{g,i} | x_g] = p_g, \quad \text{Var}(Q_g | x_g) = \frac{v_g}{K}. \quad (26)$$

Since $Z_{g,i}$ is included in Q_g , their conditional covariance is

$$\text{Cov}(Q_g, Z_{g,i} | x_g) = \frac{v_g}{K}. \quad (27)$$

Linearity of expectation gives $\mathbb{E}[T_{g,i} | x_g] = p_g$. Expanding the variance yields

$$\text{Var}(T_{g,i} | x_g) = \lambda^2 \frac{v_g}{K} + (1 - \lambda)^2 v_g + 2\lambda(1 - \lambda) \frac{v_g}{K} \quad (28)$$

$$= v_g \left[\frac{1}{K} + (1 - \lambda)^2 \left(1 - \frac{1}{K} \right) \right]. \quad (29)$$

For a fixed prediction a ,

$$\mathbb{E}[(a - T_{g,i})^2 | x_g] = (a - p_g)^2 + \text{Var}(T_{g,i} | x_g). \quad (30)$$

The variance is independent of a , so the unique minimizer is $a = p_g$. $\square$

Connection to proper scoring rules. Let $\ell(a, z)$ be a binary strictly proper scoring loss, so that $p\ell(a, 1) + (1 - p)\ell(a, 0)$ is uniquely minimized at $a = p$. Extend it to a target $t \in [0, 1]$ by $\ell(a, t) = t\ell(a, 1) + (1 - t)\ell(a, 0)$. For any fixed a for which these losses are finite, unbiasedness implies

$$\mathbb{E}[\ell(a, T_{g,i}) | x_g] = p_g \ell(a, 1) + (1 - p_g) \ell(a, 0). \quad (31)$$

Thus the mixture, the group average, and an individual correctness label induce the same expected scoring loss for a prediction based on x_g . For Brier loss, $\ell(a, t) = (a - t)^2 + t(1 - t)$; the second term does not depend on a , giving the minimizer stated in the proposition. This argument also applies to binary logarithmic loss through the same extension (Gneiting & Raftery, 2007).

Effect of the rollout count and mixing weight. At $\lambda = 0$, the target is the binary label with variance v_g ; at $\lambda = 1$, it is the group average with variance v_g/K . For $0 < \lambda < 1$, its variance lies between these values and decreases with K . For a fixed $\lambda < 1$, the variance approaches $(1 - \lambda)^2 v_g$ as K increases, reflecting the retained contribution of the individual outcome. Within any observed group, $T_{g,i} - T_{g,j} = (1 - \lambda)(Z_{g,i} - Z_{g,j})$. The mixing weight therefore controls the tradeoff between lower target variance and separation of correct and incorrect responses.

Conditioning on the generated response. The proposition concerns the success probability p_g before observing a particular response. To make the effect of conditioning explicit, let h_i include x_g and information from response i , and define $p_{h_i} = \Pr(Z_{g,i} = 1 | h_i)$. If the other rollouts are independent of h_i conditional on x_g , then

$$\mathbb{E}[T_{g,i} | h_i] = \left(1 - \lambda + \frac{\lambda}{K} \right) p_{h_i} + \lambda \left(1 - \frac{1}{K} \right) p_g. \quad (32)$$

The target thus shrinks the response-conditional probability toward the prompt’s success probability. The proposition justifies the mean and variance of this supervision signal under the original rollout distribution; it does not establish response-conditional calibration after target balancing and token-level SFT. Those effects are evaluated empirically.

D IMPLEMENTATION DETAILS

D.1 TRAINING PIPELINES

All runs use Qwen3-8B as the base model. We randomly hold out 2,000 examples from DeepScaleR-Preview as a test split (DeepScaleR-Eval) and use all remaining examples as the training split (DeepScaleR-Train), which is used to construct the SFT data for CalibSFT and the downstream RL training data.

CalibSFT. To construct CalibSFT targets, we sample 10 rollouts per prompt from DeepScaleR-Train with temperature 1.0 and top- p 0.95. We compute its group success rate, and combine group-level and rollout-level correctness with $\lambda = 0.5$. We use target-balanced sampling and the selective completion mask with $\beta = 0.5$: correct rollouts supervise both the reasoning/answer and confidence segments, whereas incorrect rollouts supervise only the confidence segment. Generally, we train CalibSFT for 180 steps with a learning rate of 2×10^{-6} and 10 warm-up steps.

SFT data construction. We compute each prompt’s success rate q_g from all sampled responses before filtering. Invalid or low-quality responses are then removed from the SFT data, with q_g kept fixed. For each retained response, we replace the confidence suffix with its target while preserving the sampled reasoning and answer. In the main probability format, the target is rendered to two decimal places in the trailing `<confidence>` tag. For balanced sampling, we group examples by the numeric target rounded to six decimal places and sample the same number from every nonempty group, using the smallest group size as the default count. We apply this rule either through a per-epoch sampler or when constructing an offline balanced split.

Confidence-aware RL. For downstream RL, each prompt group contains 8 sampled responses, with each response capped at 8,192 tokens. We train with batches of 64 prompt groups for 160 optimization steps. For methods using DAPO dynamic sampling, this corresponds to approximately 600 rollout-generation batches. Dynamic sampling retains only groups with mixed correctness outcomes and discards groups whose responses are all correct or all incorrect. To accelerate training for these methods, we remove questions for which all sampled responses from the base model are correct before RL training begins. Each method retains its own reward and credit assignment.

- **Direct RL baselines.** RLVR, RLCR (Damani et al., 2026), CoCA (Li et al., 2026), and DCPO (Ma et al., 2026) are trained directly from the Qwen3-8B base checkpoint. They share the DAPO training recipe but retain their respective correctness, calibration-target, and credit-assignment objectives.
- **ReDoubt.** ReDoubt is initialized from the trained RLVR checkpoint and then optimized with its clipped log-score confidence objective (Bani-Harouni et al., 2026). When incorporated with CalibSFT, we first train CalibSFT on the RLVR checkpoint and then continue with ReDoubt RL.
- **C3RL.** C3RL (Yang et al., 2026) combines correctness, calibration, and reference-accuracy rewards. C3RL uses the digit confidence format, expressing confidence as an integer from 0 to 9 and mapping it to $[0, 1]$ by dividing by 9. The CalibSFT initialization for C3RL is also trained in this format. We use the base model’s rollout success rates to label each training question as all correct, partially correct, or all incorrect. The reference reward encourages correct answers to previously all-incorrect questions and penalizes incorrect answers to previously all-correct questions. To retain these questions and their reference-reward signals, we disable DAPO dynamic sampling and all-correct filtering, using the full RL training split for C3RL both with and without CalibSFT initialization.

D.2 EVALUATION PROTOCOL

For the main results, we use temperature 0.6 and sample four rollouts per question on non-AIME mathematical and OOD benchmarks and 32 rollouts per question on AIME. Pass@1 is the mean correctness over these sampled rollouts. The supervision and target ablations retain their recorded evaluation settings. Confidence is parsed from the final probability field and normalized to $[0, 1]$. We use 10 equal-width bins for calculating ECE. Unless otherwise noted, RL results use the final

checkpoint after 160 optimization steps, while SFT-only results use the checkpoint at the end of their respective SFT runs.

E DETAILED EXPERIMENTAL RESULTS

E.1 PER-DATASET MAIN RESULTS

Table 1 reports split-level averages. The following tables expose the same metrics for every benchmark; the final cell in each split is the unweighted average over its datasets.

E.2 ABLATION OF SFT CONFIDENCE TARGETS

To evaluate the advantage of confidence targets derived from a proper scoring rule, we compare CalibSFT with targets based on group-level correctness (CSFT) (Jang et al., 2025), token probabilities (UaIT) (Liu et al., 2024), semantic consistency (SaySelf) (Xu et al., 2024), and rubric-based LLM judgments (LoVeC) (Zhang et al., 2026). For a fair comparison, we construct the training datasets from a shared response pool. SaySelf uses its correct responses, following its SFT setup. We sample questions uniformly across empirical success rates so that the comparison covers both easier and harder questions.

For each question in DeepScaleR-Train, we use 10 responses generated by Qwen3-8B. With $z_{g,i} \in \{0, 1\}$ denoting the response correctness, the empirical success rate is $q_g = \frac{1}{10} \sum_{i=1}^{10} z_{g,i}$.

We select 300 questions at each $q_g \in \{0.1, 0.2, \dots, 0.9\}$ and retain one correct and one incorrect response per question. Finally, We have 5,400 responses from 2,700 questions, with equal numbers of correct and incorrect responses at each success rate.

On this shared training subset, we construct the five confidence targets as follows:

- **CSFT.** Both responses receive the group success rate, $c_{g,i}^* = q_g$.
- **UaIT.** The target is $c_{g,i}^* = \exp(-\ell_{g,i})$, where $\ell_{g,i}$ is the mean negative log-likelihood of the sampled completion tokens under the generating model.
- **SaySelf.** SaySelf uses an embedding model to group semantically similar responses and derives confidence from cluster frequency. We embed the ten responses to each question using Qwen3-Embedding-8B and cluster them with a cosine similarity threshold of 0.98. Since SaySelf uses only correct responses for SFT, we retain the 2,700 selected correct responses and assign each the frequency of its cluster, $c_{g,i}^* = |C_{g,i}|/10$.
- **LoVeC.** LoVeC uses an LLM judge to assess the factuality of generated statements and supervise their confidence estimates. Our rubric assigns 10 to correct answers with sound reasoning, 8–9 to correct answers with minor reasoning gaps, 5–7 to partially correct or ambiguous responses, 2–4 to meaningful but incorrect attempts, and 0–1 to invalid or clearly wrong responses. We normalize its score $s_{g,i} \in \{0, \dots, 10\}$ to obtain $c_{g,i}^* = s_{g,i}/10$.
- **CalibSFT.** Follow our official implementation, the target combines group success rate with response correctness, $c_{g,i}^* = 0.5q_g + 0.5z_{g,i}$.

All variants initialize from Qwen3-8B and use the same supervision objective. Reasoning and answer tokens receive supervision on correct responses, while confidence tokens receive supervision on all responses. We train for 80 steps (40 for SaySelf) with a batch size of 256 and a learning rate of 2×10^{-6} . We report the proportion of training targets in each of 10 equal-width confidence intervals for all five methods in Table 10.

Table 10: Training-target distribution (%) over ten equal-width bins in the confidence target ablation.

Interval	CSFT	UaIT	SaySelf	LoVeC	CalibSFT
[0.0, 0.1]	11.11	0.00	22.07	27.33	11.11
(0.1, 0.2]	11.11	0.00	14.41	3.78	11.11
(0.2, 0.3]	11.11	0.00	11.93	8.17	11.11
(0.3, 0.4]	11.11	0.00	10.22	0.20	11.11
(0.4, 0.5]	11.11	0.00	9.04	1.37	5.56
(0.5, 0.6]	11.11	0.00	8.26	0.63	11.11
(0.6, 0.7]	11.11	6.48	8.89	0.22	11.11
(0.7, 0.8]	11.11	56.06	6.63	1.54	11.11
(0.8, 0.9]	11.11	35.15	5.56	15.80	11.11
(0.9, 1.0]	0.00	2.31	3.00	40.96	5.56

Table 8: Per-dataset results on the in-domain mathematical benchmarks. Base, RLVR, and standalone CalibSFT use 6,144 tokens; downstream RL models use 16,384 tokens. Temperature is 0.6, with four rollouts per non-AIME question and 32 per AIME question.

(a) DeepScaleR-Eval					(b) MATH-500					(c) MinervaMath				
Method	P@1	AUROC	Brier	ECE	Method	P@1	AUROC	Brier	ECE	Method	P@1	AUROC	Brier	ECE
Base	60.84	66.31	34.09	34.06	Base	83.90	63.80	14.78	13.87	Base	39.89	63.26	53.79	55.28
RLVR	70.15	68.11	26.94	26.37	RLVR	89.15	67.69	10.13	8.76	RLVR	37.96	60.92	57.18	58.40
CalibSFT	60.88	73.76	23.60	17.57	CalibSFT	83.65	82.11	14.70	11.53	CalibSFT	39.89	57.50	30.33	24.20
RLCR	71.24	75.33	16.65	5.14	RLCR	90.50	76.39	7.54	7.43	RLCR	36.67	66.67	35.91	38.12
+ CalibSFT	72.47	79.34	15.33	4.58	+ CalibSFT	90.15	89.79	6.56	3.69	+ CalibSFT	37.04	76.29	26.81	26.33
ReDoubt	71.53	75.92	17.13	10.28	ReDoubt	90.05	74.85	9.82	15.01	ReDoubt	37.59	66.65	27.92	25.37
+ CalibSFT	72.36	78.06	16.06	6.67	+ CalibSFT	89.45	89.64	9.60	16.74	+ CalibSFT	36.21	83.73	21.06	22.86
CoCA	70.16	73.17	17.73	8.51	CoCA	89.85	72.69	10.64	18.18	CoCA	35.20	65.28	32.43	33.26
+ CalibSFT	71.45	78.71	16.41	7.92	+ CalibSFT	89.85	87.29	9.09	15.49	+ CalibSFT	36.95	78.62	23.47	24.42
DCPO	71.45	71.74	17.30	6.50	DCPO	89.30	70.77	7.90	6.27	DCPO	34.93	64.19	41.78	44.59
+ CalibSFT	70.46	77.71	16.06	5.95	+ CalibSFT	90.20	84.89	7.68	7.70	+ CalibSFT	36.86	74.73	23.33	20.25
C3RL	70.65	62.87	24.64	25.05	C3RL	89.40	60.22	9.42	9.75	C3RL	37.22	56.00	58.42	59.49
+ CalibSFT	71.05	71.91	19.93	14.72	+ CalibSFT	87.70	84.76	10.37	9.36	+ CalibSFT	35.85	77.64	20.14	16.95
(d) OlympiadBench					(e) GSM8K					(f) AIME2024				
Method	P@1	AUROC	Brier	ECE	Method	P@1	AUROC	Brier	ECE	Method	P@1	AUROC	Brier	ECE
Base	54.41	69.82	38.93	40.13	Base	94.35	67.82	5.20	3.39	Base	48.23	81.93	39.61	42.50
RLVR	56.31	72.84	39.35	39.98	RLVR	94.64	67.97	5.06	3.05	RLVR	52.92	82.05	39.49	41.15
CalibSFT	51.67	79.33	22.48	20.62	CalibSFT	94.31	70.24	17.90	20.78	CalibSFT	45.42	89.96	11.93	6.03
RLCR	57.83	81.78	19.04	12.50	RLCR	94.81	78.88	4.77	5.58	RLCR	59.48	93.59	13.17	19.70
+ CalibSFT	58.46	89.71	16.09	13.25	+ CalibSFT	94.67	82.60	4.64	2.72	+ CalibSFT	64.17	92.42	10.97	9.79
ReDoubt	57.57	80.67	18.18	8.36	ReDoubt	94.96	78.40	7.07	16.27	ReDoubt	55.62	91.94	15.40	22.77
+ CalibSFT	58.35	91.01	14.96	16.88	+ CalibSFT	94.69	80.28	7.50	17.62	+ CalibSFT	61.56	93.81	11.89	14.37
CoCA	56.38	80.19	20.14	14.63	CoCA	94.58	69.00	8.94	21.02	CoCA	53.85	92.22	16.27	21.37
+ CalibSFT	56.19	90.47	18.09	20.16	+ CalibSFT	94.56	71.85	8.26	18.10	+ CalibSFT	58.54	95.41	16.34	23.56
DCPO	57.05	78.43	21.85	17.75	DCPO	94.48	68.34	5.50	9.13	DCPO	57.92	88.84	12.32	8.47
+ CalibSFT	57.46	88.53	15.72	12.14	+ CalibSFT	94.45	73.40	5.80	7.48	+ CalibSFT	58.33	93.89	11.02	11.85
C3RL	57.68	62.66	35.85	37.52	C3RL	94.37	57.25	5.33	5.24	C3RL	59.69	80.93	25.25	30.19
+ CalibSFT	56.64	84.53	17.91	18.60	+ CalibSFT	94.29	64.91	5.26	5.08	+ CalibSFT	64.27	85.24	14.04	12.99
(g) AIME2025					(h) AIME2026					(i) Overall (ID)				
Method	P@1	AUROC	Brier	ECE	Method	P@1	AUROC	Brier	ECE	Method	P@1	AUROC	Brier	ECE
Base	35.00	82.75	48.25	52.81	Base	42.50	76.08	45.13	48.07	Base	57.39	71.47	34.97	36.26
RLVR	43.44	84.95	45.87	49.00	RLVR	48.23	80.70	43.82	45.92	RLVR	61.60	73.16	33.48	34.08
CalibSFT	34.27	92.48	9.95	8.02	CalibSFT	39.48	88.65	13.22	9.79	CalibSFT	56.20	79.25	18.01	14.82
RLCR	43.33	95.53	13.60	23.06	RLCR	47.81	94.10	14.08	23.19	RLCR	62.71	82.78	15.60	16.84
+ CalibSFT	48.02	92.40	13.87	14.28	+ CalibSFT	52.60	92.48	13.57	12.04	+ CalibSFT	64.70	86.88	13.48	10.84
ReDoubt	44.90	94.22	15.19	26.51	ReDoubt	52.40	90.93	16.78	26.20	ReDoubt	63.08	81.70	15.94	18.85
+ CalibSFT	51.77	96.32	11.28	16.79	+ CalibSFT	55.52	93.82	11.66	17.69	+ CalibSFT	64.99	88.33	13.00	16.20
CoCA	40.31	95.39	17.04	29.31	CoCA	48.54	96.74	14.84	27.72	CoCA	61.11	80.59	17.25	21.75
+ CalibSFT	45.00	94.21	19.99	26.43	+ CalibSFT	48.65	93.96	19.52	25.63	+ CalibSFT	62.65	86.31	16.40	20.21
DCPO	41.15	91.19	13.33	14.30	DCPO	49.58	91.78	11.04	10.41	DCPO	61.98	78.16	16.38	14.68
+ CalibSFT	41.25	95.07	13.81	17.73	+ CalibSFT	50.21	93.07	12.87	14.52	+ CalibSFT	62.40	85.16	13.29	12.20
C3RL	43.96	79.01	36.25	42.69	C3RL	52.81	84.29	27.21	33.72	C3RL	63.22	67.90	27.80	30.45
+ CalibSFT	43.12	86.74	21.86	25.49	+ CalibSFT	53.65	91.31	15.69	18.30	+ CalibSFT	63.32	80.88	15.65	15.18

Table 9: Per-dataset results on the revised eight-dataset out-of-domain suite. Each question has four sampled responses at temperature 0.6. WebQuestions contains only examples with a single gold answer.

(a) HotpotQA					(b) TriviaQA					(c) DROP				
Method	P@1	AUROC	Brier	ECE	Method	P@1	AUROC	Brier	ECE	Method	P@1	AUROC	Brier	ECE
Base	64.68	57.68	32.82	32.16	Base	55.24	75.80	35.68	36.93	Base	80.36	59.13	18.22	17.03
RLVR	64.58	54.16	33.10	32.17	RLVR	54.97	76.01	40.59	41.40	RLVR	83.53	59.32	15.34	13.31
CalibSFT	65.70	54.18	29.38	20.57	CalibSFT	55.21	65.05	28.44	20.98	CalibSFT	79.92	59.38	31.60	28.93
RLCR	65.80	59.10	25.97	19.77	RLCR	55.24	81.66	23.84	23.30	RLCR	81.44	63.34	14.72	5.43
+ CalibSFT	66.27	59.23	24.25	10.85	+ CalibSFT	55.25	81.53	18.89	10.71	+ CalibSFT	82.32	74.44	14.99	10.75
ReDoubt	64.83	60.64	23.52	12.11	ReDoubt	55.69	82.96	18.67	13.58	ReDoubt	83.63	63.27	14.76	10.80
+ CalibSFT	64.68	63.08	21.70	1.37	+ CalibSFT	56.14	83.92	19.27	15.58	+ CalibSFT	82.53	75.90	15.60	17.36
CoCA	64.90	56.43	22.77	6.02	CoCA	54.77	77.62	21.63	12.97	CoCA	82.47	61.58	15.03	11.56
+ CalibSFT	66.10	61.57	22.92	11.39	+ CalibSFT	54.89	81.85	20.98	13.16	+ CalibSFT	82.50	73.91	17.95	21.99
DCPO	66.15	57.88	25.33	18.77	DCPO	54.55	79.86	25.93	23.91	DCPO	83.45	59.34	13.25	2.30
+ CalibSFT	65.97	60.70	24.43	14.92	+ CalibSFT	55.12	82.56	17.59	8.25	+ CalibSFT	80.42	75.92	16.54	13.64
C3RL	65.38	52.47	33.39	33.12	C3RL	53.76	73.01	35.61	38.53	C3RL	82.65	54.56	16.37	16.07
+ CalibSFT	66.50	58.94	28.45	19.78	+ CalibSFT	54.11	71.36	22.75	12.50	+ CalibSFT	77.44	64.78	23.65	22.03
(d) MuSiQue					(e) LiveBenchReasoning					(f) NQOpen				
Method	P@1	AUROC	Brier	ECE	Method	P@1	AUROC	Brier	ECE	Method	P@1	AUROC	Brier	ECE
Base	45.84	62.90	41.91	42.33	Base	52.88	61.81	40.53	40.59	Base	21.12	66.31	66.04	70.59
RLVR	47.18	58.12	46.19	46.75	RLVR	29.25	60.32	61.75	64.59	RLVR	24.13	67.07	68.92	71.55
CalibSFT	47.19	62.25	26.51	13.95	CalibSFT	53.00	67.67	25.40	17.74	CalibSFT	21.64	59.59	27.33	29.16
RLCR	47.77	66.21	25.68	15.35	RLCR	31.75	68.00	28.49	29.09	RLCR	23.32	70.48	43.84	52.43
+ CalibSFT	49.97	67.87	23.22	6.13	+ CalibSFT	42.25	68.41	21.94	7.75	+ CalibSFT	23.80	72.24	26.44	30.89
ReDoubt	46.90	66.37	24.07	11.21	ReDoubt	25.75	68.44	23.51	24.86	ReDoubt	24.67	70.28	32.15	39.19
+ CalibSFT	48.52	70.04	22.42	5.52	+ CalibSFT	42.25	80.49	19.95	12.38	+ CalibSFT	24.10	70.60	26.80	31.99
CoCA	46.69	67.03	23.69	10.18	CoCA	32.12	66.75	25.79	23.64	CoCA	22.76	66.59	35.28	43.26
+ CalibSFT	48.76	66.88	23.56	7.72	+ CalibSFT	41.50	68.03	23.24	7.96	+ CalibSFT	23.72	70.77	24.86	28.69
DCPO	47.34	65.48	28.28	20.53	DCPO	31.13	67.14	32.10	33.12	DCPO	23.69	69.80	46.27	53.69
+ CalibSFT	48.96	70.49	22.75	8.55	+ CalibSFT	37.75	71.77	21.12	8.83	+ CalibSFT	23.58	73.13	22.65	25.97
C3RL	45.55	63.39	39.88	40.04	C3RL	33.00	57.18	53.55	55.46	C3RL	22.43	64.77	63.63	68.49
+ CalibSFT	47.47	66.39	24.46	11.87	+ CalibSFT	44.62	53.33	28.11	22.78	+ CalibSFT	22.08	61.18	22.25	17.54
(g) PopQA					(h) WebQuestions					(i) Overall (OOD)				
Method	P@1	AUROC	Brier	ECE	Method	P@1	AUROC	Brier	ECE	Method	P@1	AUROC	Brier	ECE
Base	20.85	74.94	57.89	64.10	Base	21.88	67.10	68.16	71.80	Base	45.36	65.71	45.15	46.94
RLVR	21.37	72.78	68.14	71.96	RLVR	26.02	65.76	68.29	70.41	RLVR	43.88	64.19	50.29	51.52
CalibSFT	21.00	63.88	24.80	26.68	CalibSFT	23.59	59.02	32.52	35.25	CalibSFT	45.91	61.38	28.25	24.16
RLCR	21.29	76.88	37.71	47.90	RLCR	25.07	69.40	48.86	56.11	RLCR	43.96	69.38	31.14	31.17
+ CalibSFT	21.16	76.55	24.76	32.22	+ CalibSFT	24.63	70.20	31.41	36.33	+ CalibSFT	45.71	71.31	23.24	18.20
ReDoubt	21.37	76.83	26.82	36.23	ReDoubt	26.84	70.22	37.08	43.81	ReDoubt	43.71	69.88	25.07	23.98
+ CalibSFT	21.25	77.87	24.54	32.32	+ CalibSFT	24.91	69.57	29.11	34.59	+ CalibSFT	45.55	73.93	22.42	18.89
CoCA	21.06	73.92	30.84	39.88	CoCA	24.65	64.40	38.35	45.46	CoCA	43.68	66.79	26.67	24.12
+ CalibSFT	20.88	76.99	24.12	31.76	+ CalibSFT	26.06	69.71	25.93	28.47	+ CalibSFT	45.55	71.21	22.95	18.89
DCPO	20.99	77.02	41.10	49.96	DCPO	24.59	68.48	50.52	57.04	DCPO	43.99	68.13	32.85	32.41
+ CalibSFT	21.18	76.52	21.29	28.02	+ CalibSFT	25.61	70.87	27.45	30.32	+ CalibSFT	44.82	72.75	21.73	17.31
C3RL	20.81	71.98	54.60	61.92	C3RL	23.94	60.19	67.27	70.35	C3RL	43.44	62.19	45.54	48.00
+ CalibSFT	20.66	66.98	16.56	12.95	+ CalibSFT	23.94	61.29	24.87	19.55	+ CalibSFT	44.60	63.03	23.89	17.38

E.3 TRAINING STABILITY

[TODO:] Show training curves and report the mean and standard deviation of evaluation metrics across training seeds.

E.4 ADDITIONAL METRICS

[TODO:] Report the additional calibration and selective-prediction metrics, including MacroCE, adaptive ECE, PCE, and AURC/EAURC, together with their definitions and aggregation protocols. Include ECE sensitivity to 10, 15, and 20 bins, and plots comparing calibration metrics with AUROC and Pass@1.

E.5 GENERALIZATION TO ADDITIONAL MODELS

[TODO:] Evaluate CalibSFT and representative downstream RL objectives with OLMo and Gemma, using the same ID/OOD protocol as the main experiment.

E.6 CONFIDENCE VISUALIZATION

[TODO:] Show reliability diagrams and confidence distributions for models with and without CalibSFT initialization.

E.7 CONFIDENCE DISTRIBUTION ANALYSIS

[TODO:] Report Top-1 and Top-3 confidence mass and confidence entropy, with per-dataset results and ID/OOD averages.

E.8 TRAINING EFFICIENCY

[TODO:] Compare CalibSFT and confidence-aware RL using wall-clock time, GPU-hours, optimization steps, and generated rollout/token counts.

E.9 QUALITATIVE CASE STUDIES

[TODO:] Include representative high-confidence errors, low-confidence correct answers, confidence corrections preserved after RL, and failure cases.